

Agentic Ontology of Work

Objective → Intent → Plan → Task → Action (by an Actor, with a Skill) → Result → Outcome → Feedback → Memory. Governed by Policy, Confidence, Assurance Level, and the Guardian.

Perception

Records the work in progress and what is observed about it.

Work Item The unit of business work that the rest of the ontology is about: a claim, an application, an order, a case, a ticket.

Signal A raw, time-stamped trace emitted by a system, a person, or the environment: an event, a log line, a user-interface interaction, a document arriving, a state change, a metric sample.

Observation An interpreted fact about a Work Item, Actor, or Action, derived from one or more Signals and carrying a Confidence.

Cognition

Defines what should happen and the limits that apply.

Objective A business-level goal expressed in strategic or operational terms.

Intent A structured, actionable goal derived from an Objective and interpretable by agents, orchestrators, and people.

Context The structured situational information that gives relevance to an Intent.

Policy Declarative constraints defining permissible, required, or restricted behavior.

Plan An ordered or partially ordered set of Tasks intended to fulfill an Intent, produced by an Orchestrator or other Actor.

Execution

Records who performed the work and what it changed.

Task A unit of assigned work within a Plan, performed by an Actor filling a Role.

Actor Anything that can be assigned Tasks, perform Actions, and be held to account for them: an Agent, a person, an Orchestrator, or a Guardian.

Agent A software Actor that interprets Intents and Tasks and acts on them using Skills, within the autonomy it has been cleared for.

Human Actor A person who performs, approves, reviews, or oversees work.

Orchestrator A coordinating Actor that turns Intents into Plans, assigns Tasks to Actors, and manages sequencing and exceptions.

Role A named set of responsibilities and authority that an Actor can fill, such as "claims adjuster" or "eligibility reviewer".

Skill A reusable capability that an Actor can invoke.

Action A single execution event in which an Actor invokes a Skill, in service of a Task, at a point in time.

Result The immediate, atomic effect or output of an Action.

Assurance

Governs autonomy, measures outcomes, and records feedback.

Guardian The oversight Actor that enforces Policy and Assurance Levels, monitors behavior, and escalates to people when needed.

Confidence A model-, rule-, or human-derived measure of certainty about an Observation, a Result, or an Action.

Assurance Level A governance-set level of permitted autonomy, with the human oversight and evidence it requires.

Outcome The aggregated business impact of one or more Results, measured against an Objective.

Feedback Structured information, from people or systems, used to improve future behavior.

Memory Persistent knowledge accumulated from experience, used to inform future Context.

Operational Memory Short-lived, execution-oriented memory scoped to a session, Task, or Work Item.

Knowledge Base Long-lived memory of generalized patterns, rules, and precedents.

Assurance Levels. The effective level is the lowest of the levels set by the Intent, Task, Skill, and Agent. Below the confidence threshold, the Agent escalates to a person.

AL0 Manual

Observes and prepares; takes no Action that changes a Work Item.

AL1 Assisted

Recommends, drafts, and prepares.

AL2 Delegated

Decides and acts within Policy.

AL3 Supervised

Acts, and must be able to reverse or compensate its Actions.

AL4 Autonomous

Acts without routine human review; the Guardian can suspend it at any time.